


Small-K Prompt Sensitivity: Bias and Directional Guidance


1,000 Monte Carlo replicates per condition

(A) Relative bias of $\hat{\sigma}_p^2$ vs K



σ_p^2 Large (0.10) Moderate (0.04)

(B) Directional accuracy near decision boundary



K = 10 K = 3 K = 5
K = 2 K = 4 K = 7