



(A) Aggregation improves accuracy

Each step averages over an additional pipeline factor




(B) TEE identifies unreliable items

Items split into tertiles by cross-config SD



(C) TEE reliability predicts accuracy

Cross-config variance predicts classification accuracy



(D) Pipeline choice creates a 18 pp accuracy spread

45 configs (5 prompts $\times$ 3 judges $\times$ 3 temps)

